Fix or *fixed*? Repository norms and subject omission in Git commits

English null subjects occur in a small set of registers, including diaries (Haegeman 1997, 2013; Weir 2012). Git commit subjects offer a new test case. Developer guides usually call strings such as *Fix parser* imperatives. The label supplies an understood YOU and explains the base form at once. Yet commits report work that has already been done, and some projects require *Fixed parser* instead. We ask whether verbal form follows the discourse function of the commit or a local writing norm. Fourteen public repositories show a sharp answer: thirteen overwhelmingly use the base, while Django uses active past tense. Both patterns omit the author. The single-community contrast separates subject omission from imperative morphology and exposes a role for repository convention in a tightly constrained technical register.

Data and method. We collected the latest 1,000 non-merge commit subjects from each of 14 major public repositories (14,000 subjects total) on November 5, 2026. The sample includes programming languages and web frameworks as well as data-science projects. Bracketed review tags and one repository scope prefix (e.g., *net/http:*) were treated as metadata outside the clause. Automatic merges and backports were excluded. To avoid relying on a general parser trained on edited prose, we used a conservative, auditable criterion: the first clausal token had to instantiate one of 40 hand-specified English paradigms whose spelling distinguishes BASE, PAST, 3SG, and PROGRESSIVE (e.g., *fix/fixed/fixes/fixing*). This retained 8,291 subjects (59.2%). The released script reproduces each coding decision from the public histories.

Results. Of the retained subjects, 7,484 began with a base form, 752 with a past form, 23 with third-person singular, and 32 with a progressive. Table 1 reports the base–past comparison. Thirteen repositories selected the base form in 97.4–100% of tokens; Django selected the past in all 685 base–past tokens. As a descriptive measure, repository identity accounts for most of the base–past association, $\chi^2(13) = 7442.9$, $p < .001$, Cramer's $V = .951$ ($N = 8{,}236$). We do not treat token-level significance as population-level inference because all past-dominant data come from Django. The lexical control is more informative. Twenty-three lemmas occur in Django and in at least one comparison repository. Django has 685 past and zero base spellings for those lemmas; the other repositories have 6,931 base and 62 past spellings (99.1% base). For *fix* alone, Django contributes 523 *fixed* tokens, while the other repositories contribute 2,213 *fix* and 14 *fixed* tokens. A fixed-seed manual check of 80 base and 80 past items found a verbal use in 159/160 cases. Past-form strings such as (1b) contain an overt direct object, which rules out a reduced passive reading.

(1) a. *Fix host mismatch in the test.*
    b. *Fixed #37388 – Improved unsafe user input examples.*

Table 1: Clause-initial base and past forms by repository.

Repository	Base	Past	Repository	Base	Past
CPython	564	1	Requests	643	16
React	496	3	NumPy	460	6
Django	0	685	pandas	281	2
Go	467	2	FastAPI	879	0
Node.js	455	0	Vue	887	2
PyTorch	433	3	Svelte	683	4
Flask	639	12	HTTPX	597	16

Django's public contributor guide requires past tense and gives *Fixed Unicode bug* as its model (Django Software Foundation 2026). Guides for several comparison projects prescribe the imperative base. The corpus distribution tracks these explicit repository-level norms.

Analysis. English imperatives lack past-tense morphology, so the Django pattern requires written-register subject omission. The missing argument is the responsible developer or team, identified in the author field beside every commit. The clause is a root, and its gap is leftmost after non-clausal scope labels have been removed. These are the central structural restrictions on diary drop (Haegeman 2013). A base-form commit still reports a completed change; its form alone does not establish directive force. We treat the two patterns as local realizations of a COMMIT-REPORT construction:

$$\left[\textsc{commit-report}\ \varnothing_{\text{author/team}}\ V_{\{\text{base,past}\}}\ XP\right].$$

The construction licenses a recoverable left-edge subject. Repository convention selects its verbal realization. Under this analysis, *imperative* remains a style-guide label; base morphology does not settle the clause's syntax or illocutionary force.

Two follow-up tests distinguish this account from a fixed genre grammar. Contributors who move between projects should adopt the destination project's tense norm. In an acceptability experiment, hiding the author field should hurt past-form null-subject reports more than versions with an overt subject. The present corpus supplies the starting observation: identical lemmas performing the same reporting task take different morphology at a repository boundary.

References

[Django Software Foundation 2026] Django Software Foundation. 2026. Committing code. `https://docs.djangoproject.com/en/dev/internals/contributing/committing-code/`.

[Haegeman 1997] Haegeman, Liliane. 1997. Register variation, truncation, and subject omission in English and French. *English Language and Linguistics* 1(2). 233–270.

[Haegeman 2013] Haegeman, Liliane. 2013. The syntax of registers: Diary subject omission and the privilege of the root. *Lingua* 130. 88–110.

[Weir 2012] Weir, Andrew. 2012. Left-edge deletion in English and subject omission in diaries. *English Language and Linguistics* 16. 105–129.